

Kubernetes & Shared Responsibility

Understanding one's scope of responsibility in a Kubernetes cloud service

Why are we here?

Kubernetes is

- Big
- Complex
- New

Overview

- Shared Responsibility Models
- Kubernetes (aka K8)
- Cloud Delivery Model
- Customer Surface Area
- Threat Matrix
- Hardening
- Q&A

Who Am I?

Ken Netzorg

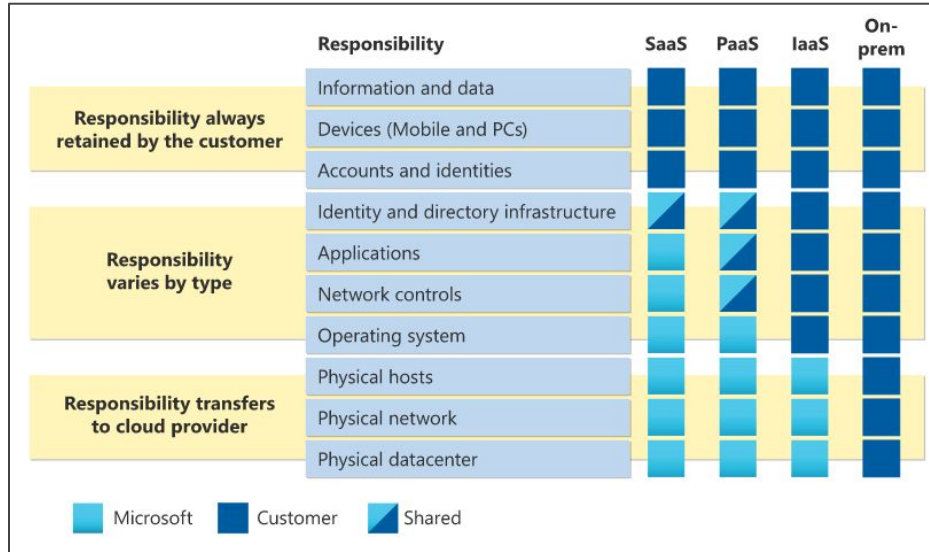
25 Years across various responsibilities ranging from network administration to DBA to C# development.

I like to solve problems, IT related.

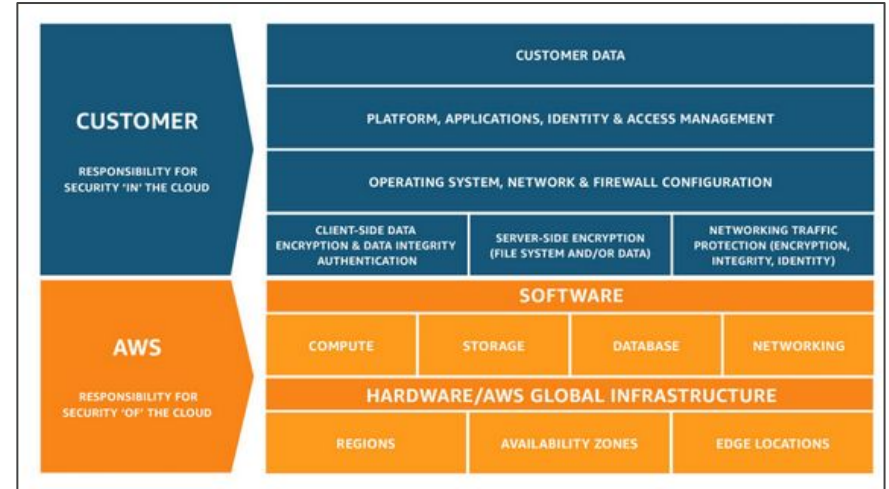
Currently the Director of Technology Operations and Security at DecisivEdge and heavily involved in Azure infrastructure, security, and DevOps.

Certs: CISA/CISSP

Shared Responsibility Models

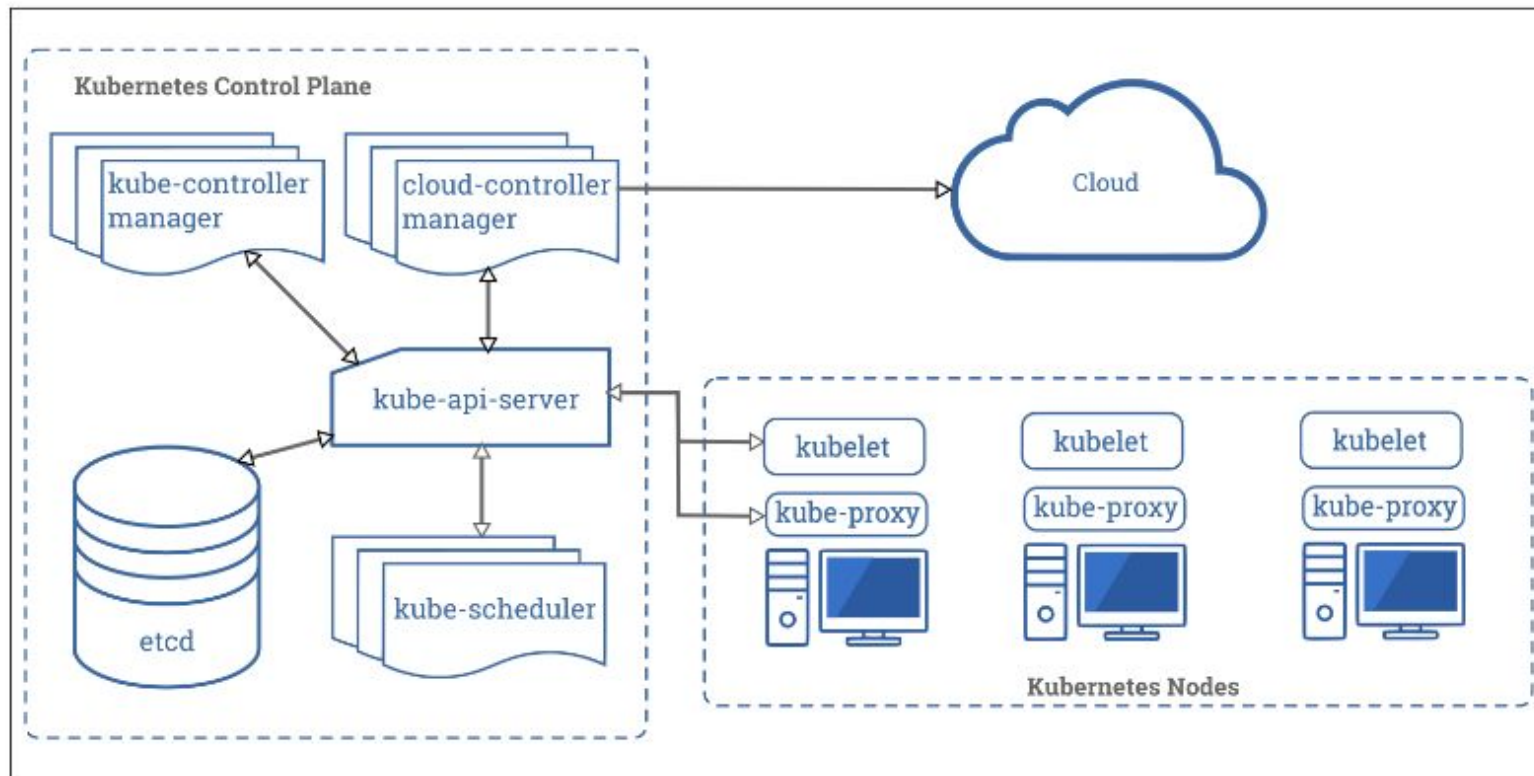


Microsoft - Azure



Amazon - AWS

Kubernetes Architecture



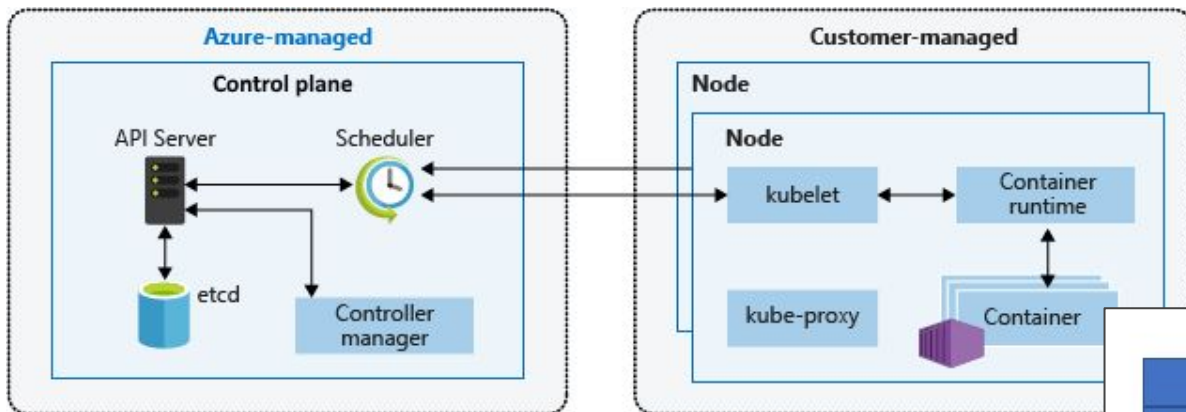
Kubernetes Architecture - Nodes



Add-Ons

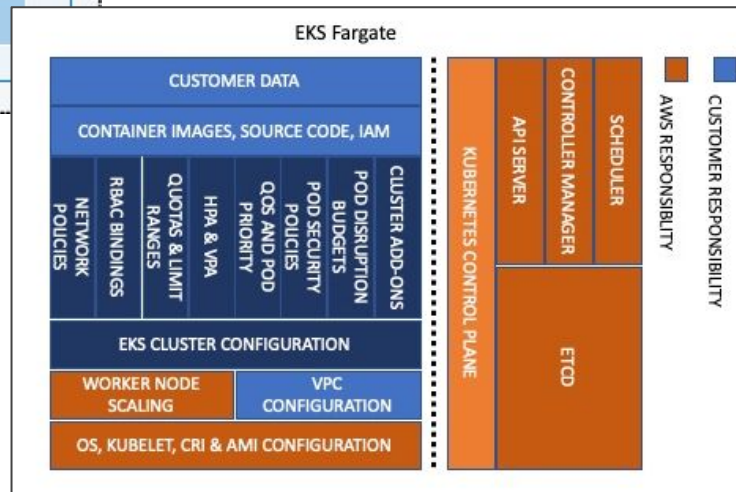
- Network Policies
- DNS
- Service Mesh
- Ingress
- Etc.

Kubernetes In the Cloud



Microsoft - Azure

Amazon - AWS





Customer Surface Area

- Application Security
 - Secure Coding
 - Patching
 - etc
- Access Management
- Transport Security
 - Encrypted
 - Segmented
- Backups/Redundancy
 - Alternate data locations
- Secret Management
- Patching



Kubernetes Threat Matrix (updated)

Initial Access	Execution	Persistence	Privilege Escalation	Defense Evasion	Credential Access	Discovery	Lateral Movement	Collection	Impact
Using Cloud credentials	Exec into container	Backdoor container	Privileged container	Clear container logs	List K8S secrets	Access the K8S API server	Access cloud resources	Images from a private registry	Data Destruction
Compromised images in registry	bash/cmd inside container	Writable hostPath mount	Cluster-admin binding	Delete K8S events	Mount service principal	Access Kubelet API	Container service account		Resource Hijacking
Kubeconfig file	New container	Kubernetes CronJob	hostPath mount	Pod / container name similarity	Access container service account	Network mapping	Cluster internal networking		Denial of service
Application vulnerability	Application exploit (RCE)	Malicious admission controller	Access cloud resources	Connect from Proxy server	Applications credentials in configuration files	Access Kubernetes dashboard	Applications credentials in configuration files		
Exposed Dashboard	SSH server running inside container				Access managed identity credential	Instance Metadata API	Writable volume mounts on the host		
Exposed sensitive interfaces	Sidecar injection				Malicious admission controller		Access Kubernetes dashboard		
							Access tiller endpoint		
							CoreDNS poisoning		
							ARP poisoning and IP spoofing		

 = New technique
 = Deprecated technique

Hardening Help

Sources are available, all is not lost

- NIST - Kubernetes Hardening Guidance (total platform)
- CIS - Total platform and specific provider guides
- Azure - “Day 2” Topics
- Amazon - EKS Best Practices
- Azure AKS Baseline Cluster (Azure centric/full solution concept)

Links to these documents will be in the final slide

Hardening Help - NIST

Contents

Kubernetes Hardening Guidance	i
Executive summary	iv
Introduction	1
Recommendations	2
Architectural overview	3
Threat model	5
Kubernetes Pod security	7
"Non-root" containers and "rootless" container engines	7
Immutable container file systems	8
Building secure container images	8
Pod Security Policies	10
Protecting Pod service account tokens	11
Hardening container engines	12
Network separation and hardening	13
Namespaces	13
Network policies	14
Resource policies	14
Control plane hardening	15
EtcD	16
Kubeconfig Files	16
Worker node segmentation	16
Encryption	17
Secrets	17
Protecting sensitive cloud infrastructure	18
Authentication and authorization	18
Authentication	19
Role-based access control	20
Log auditing	22
Logging	22
Kubernetes native audit logging configuration	24
Worker node and container logging	25
Seccomp: audit mode	26
SYSLOG	27
SIEM platforms	27
Alerting	28
Service meshes	29
Fault tolerance	30
Tools	31
Upgrading and application security practices	32

NIST main TOC

Hardening Help - CIS (Azure)

Table of Contents

Terms of Use.....	1
Overview.....	6
Intended Audience.....	6
Consensus Guidance.....	6
Typographical Conventions.....	7
Assessment Status.....	7
Profile Definitions.....	8
Acknowledgements.....	9
Recommendations.....	10
1 Master (Control Plane) Components.....	10
2 Master (Control Plane) Configuration.....	11
2.1 Logging.....	12
2.1.1 Enable audit Logs (Manual).....	12
3 Worker Nodes.....	15
3.1 Worker Node Configuration Files.....	16
3.1.1 Ensure that the kubeconfig file permissions are set to 644 or more restrictive (Manual).....	17
3.1.2 Ensure that the kubelet kubeconfig file ownership is set to root:root (Manual).....	19
3.1.3 Ensure that the kubelet configuration file has permissions set to 644 or more restrictive (Manual).....	21
3.1.4 Ensure that the kubelet configuration file ownership is set to root:root (Manual).....	23
3.2 Kubelet.....	25
3.2.1 Ensure that the --anonymous-auth argument is set to false (Manual).....	25
3.2.2 Ensure that the --authorization-mode argument is not set to AlwaysAllow (Manual).....	29
3.2.3 Ensure that the --client-ca-file argument is set as appropriate (Manual).....	32
3.2.4 Ensure that the --read-only-port is secured (Manual).....	35
3.2.5 Ensure that the --streaming-connection-idle-timeout argument is not set to 0 (Manual).....	37

3.2.6 Ensure that the --protect-kernel-defaults argument is set to true (Manual).....	40
3.2.7 Ensure that the --make-iptables-util-chains argument is set to true (Manual).....	43
3.2.8 Ensure that the --hostname-override argument is not set (Manual).....	46
3.2.9 Ensure that the --eventRecordQPS argument is set to 0 or a level which ensures appropriate event capture (Manual).....	49
3.2.10 Ensure that the --rotate-certificates argument is not set to false (Manual).....	52
3.2.11 Ensure that the RotateKubeletServerCertificate argument is set to true (Manual).....	55
4 Policies.....	58
4.1 RBAC and Service Accounts.....	59
4.1.1 Ensure that the cluster-admin role is only used where required (Manual).....	59
4.1.2 Minimize access to secrets (Manual).....	61
4.1.3 Minimize wildcard use in Roles and ClusterRoles (Manual).....	63
4.1.4 Minimize access to create pods (Manual).....	65
4.1.5 Ensure that default service accounts are not actively used. (Manual).....	67
4.1.6 Ensure that Service Account Tokens are only mounted where necessary (Manual).....	69
4.2 Pod Security Policies.....	71
4.2.1 Minimize the admission of privileged containers (Automated).....	72
4.2.2 Minimize the admission of containers wishing to share the host process ID namespace (Automated).....	74
4.2.3 Minimize the admission of containers wishing to share the host IPC namespace (Automated).....	76
4.2.4 Minimize the admission of containers wishing to share the host network namespace (Automated).....	78
4.2.5 Minimize the admission of containers with allowPrivilegeEscalation (Automated).....	80
4.2.6 Minimize the admission of root containers (Automated).....	82
4.2.7 Minimize the admission of containers with the NET_RAW capability (Automated).....	84
4.2.8 Minimize the admission of containers with added capabilities (Automated).....	86

CIS partial TOC

Hardening Help - AWS EKS



EKS Best Practices Guides



aws/aws-eks-best-practices

EKS Best Practices Guides

Guides



[Introduction](#)

Security



Home

Identity and Access
Management

Pod Security

Multi-tenancy

Detective Controls

Network Security

Data Encryption and Secrets
Management

Runtime Security

Infrastructure Security

Regulatory Compliance

Incident Response and
Forensics

Image Security

Cluster Autoscaling



Reliability



Windows Containers (beta)



Introduction



Table of contents

[Contributing](#)

Welcome to the EKS Best Practices Guides. The primary goal of this project is to offer a set of best practices for day 2 operations for Amazon EKS. We elected to publish this guidance to GitHub so we could iterate quickly, provide timely and effective recommendations for variety of concerns, and easily incorporate suggestions from the broader community.

We currently have published guides for the following topics:

- [Best Practices for Security](#)
- [Best Practices for Reliability](#)
- [Best Practices for Cluster Autoscaling](#)

In the future we will be publishing best practices guidance for performance, cost optimization, and operational excellence.

Contributing

We encourage you to contribute to these guides. If you have implemented a practice that has proven to be effective, please share it with us by opening an issue or a pull request. Similarly, if you discover an error or flaw in the guidance we've already published, please submit a PR to correct it.

EKS Best Practices Intro

Conclusion

- Kubernetes is complex
- Leveraging a cloud service provider eases the burden
- Understanding where the provider stops and you start is critical
- Help is readily available if you know where to look



Q & A

Contact

Ken Netzorg

knetzorg@gmail.com

zorg_the_blue (discord)

Copy of the slides will be posted to
my github: zorg-the-blue

<https://github.com/zorg-the-blue/talks>

Document Source Links

NIST - Kubernetes Hardening Guidance (total platform)

https://media.defense.gov/2021/Aug/03/2002820425/-1/-1/1/CTR_KUBERNETES%20HARDENING%20GUIDANCE.PDF

CIS - Total platform and specific provider guides/checklists *[Requires registration]*

<https://www.cisecurity.org/benchmark/kubernetes>

Azure - “Day 2” Topics

<https://docs.microsoft.com/en-us/azure/architecture/operator-guides/aks/day-2-operations-guide>

Amazon - EKS Best Practices

<https://aws.github.io/aws-eks-best-practices/>

Azure AKS Baseline Cluster (Azure centric/full solution concept)

<https://github.com/mspnp/aks-secure-baseline>